

I certify that this is all my own original work. If I took any parts from elsewhere, then they were non-essential parts of the assignment, and they are clearly attributed in my submission. I will show that I agree to this honour code by typing "Yes": Yes

Problem 1.

- a) The algorithm computes the number of nodes in the binary tree by traversing to the leftmost node, at which recursive calls return 0 (if it does not have a right child, in which case a number 1 or greater is returned depending on the number of child nodes to its right), which is added to 1 (denoting its own count), which is then returned to a recursive call made from the parent node, which then the same recursive call to its right child before repeating this whole process up to the root node, at which the total number of nodes is returned.
- b) Divide-and-conquer: a problem is divided into subproblems, solutions to which are combined to get a solution to the original problem without discarding any of them as in decrease-and-conquer.
- c) $C(h) = C(h - 1) + 2^{h+1}$ for every $h > 0$, $C(0) = 2$
- d) The method of backward substitutions is applied to find a pattern, with which an explicit formula in terms of h can be derived:
1. $C(h) = C(h - 1) + 2 \cdot 2^h$ substitute $C(h - 1) = C(h - 2) + 2 \cdot 2^{h-1}$
 2. $C(h) = C(h - 2) + 2(2^{h-1} + 2^h)$ substitute $C(h - 2) = C(h - 3) + 2 \cdot 2^{h-2}$
 3. $C(h) = C(h - 3) + 2(2^{h-2} + 2^{h-1} + 2^h)$
 4. Now, a general formula according to the pattern above is:

$$C(h) = C(h - i) + 2(2^{h-(i-1)} + \dots + 2^{h-(i-i)})$$
 5. Plugging in h in i yields

$$C(h) = C(h - h) + 2(2^{h-(h-1)} + \dots + 2^{h-(h-h)})$$

$$C(h) = 2 + 2(2^1 + \dots + 2^h)$$

$$C(h) = 2 + 2(1 + 2^1 + \dots + 2^h) - 2$$

$$C(h) = 2(1 + 2^1 + \dots + 2^h)$$

$$C(h) = 2 \sum_{k=0}^h 2^k = 2(2^{h+1} - 1)$$
- e) $\Theta(2^n)$

Problem 2.

In all the algorithms below, all the operations will be done in-place.

- a) Description
1. Iterate over *list_triple*.
 2. In the current iteration, linearly search the passport number of the current person in *list_double*. Once it is found, remove the person from *list_double*.
 3. Repeat the step 2 for the remainder of the iteration.
 4. Once the iteration is over, return *list_double*.

The time complexity of the algorithm above is $O(mn)$.

b) Description

1. Heapsort *list_triple* by the passport number [$O(m \log m)$]

2. Iterate over *list_double* to find the current passport in *list_triple*. If it is found, remove the associated person from *list_double*; else the person is kept.
3. Repeat the step 2 until the iteration is over [$O(n \log m)$].
4. Return the remaining people in *list_double*.

The time complexity of the algorithm above is calculated as follows:

- Heapify *list_triple*: $O(m)$
- Sort the heap using siftDown() on m elements: $O(m \log m)$
- The dominating term of the two prevails: $O(m + m \log m) = O(m \log m)$
- Search in the heap, $O(\log m)$, n times and remove the person if found: $O(n \log m)$
- The dominating term of the two operations $O(m \log m)$ and $O(n \log m)$ is $O(n \log m)$ since $m \leq n$.

The requirement for the constant space complexity is satisfied as the heapsort's space complexity is $O(1)$.

c) Description

1. Build a hash table for *list_triple* and name it *hash_triple* (*HT*)
2. Iterate over *list_double* to find the current passport in *HT* and remove it from *list_double* if found or keep it otherwise
3. Once the iterations are over, return the modified version of *list_double*.

Pseudocode

ALGORITHM *FindWhoNeedBoosterShots(LD,LT)*

// Finds who have been double vaccinated but have not received third shots.
 // Input: A list *LD* of n people, represented as passport numbers, who have been double vaccinated and a list *LT* of m people, again represented as passport numbers, who have been double vaccinated and have received the third shot.
 // Output: A modified list of *LD* that contains all the passport number of the people who need the third shot

HT ← new Hashtable(LT) // hash table of m people from *LT*, with each of the passport numbers as a key and the Boolean value True as its value

FOR each passport_number pn in LD:

 IF HT contains pn:

 REMOVE pn from LD

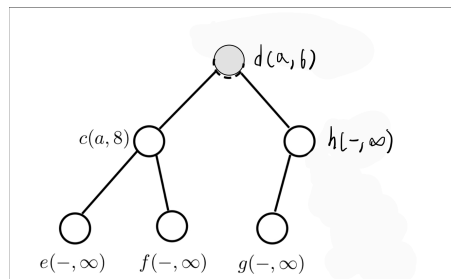
RETURN LD

Time Complexity

- $O(m)$ since the time complexity for inserting an element into a hash table is $O(1)$ on average and m insertions are required to build *HT*.
- $O(n)$ since the time complexity for searching an element in a hash table is $O(1)$ on average and n searches (iterations) are required.
- $O(m + n)$ on average where $m \leq n$ equals $O(n)$.

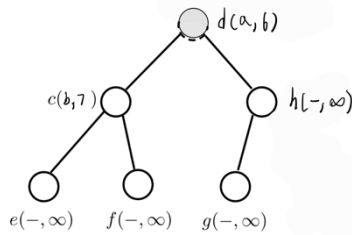
Problem 3.

a)

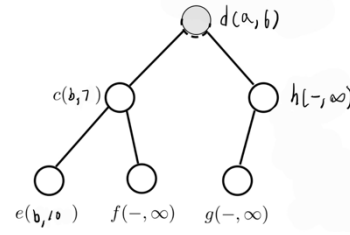


b)

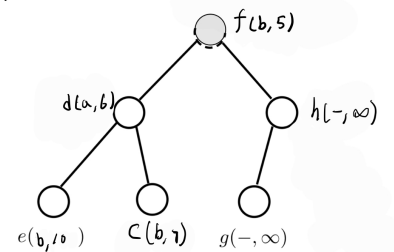
1)



2)



3)



c)

	S: vertices whose shortest paths have been known	Priority queue of remaining vertices
1	$a(a, 0)$	$b(a, 4), c(a, 8), d(a, 6), e(-, \infty), f(-, \infty), g(-, \infty), h(-, \infty)$
2	$a(a, 0), b(a, 4)$	$f(b, 5), d(a, 6), h(-, \infty), e(b, 10), c(b, 7), g(-, \infty)$
3	$a(a, 0), b(a, 4), f(b, 5)$	$d(a, 6), c(b, 7), h(f, 12), e(b, 10), g(-, \infty)$
4	$a(a, 0), b(a, 4), f(b, 5), d(a, 6)$	$c(b, 7), e(b, 10), h(f, 12), g(d, 10)$
5	same as above (SAA), $c(b, 7)$	$g(c, 9), e(b, 10), h(f, 12)$
6	SAA, $g(c, 9)$	$e(g, 9), h(g, 10)$
7	SAA, $e(g, 9)$	$h(g, 10)$
8	SAA, $h(g, 10)$	

Table 1: Complete this table for Part c).

d)

	Shortest Paths	Distances
a	$a \rightarrow a$	0
b	$a \rightarrow b$	4
c	$a \rightarrow b \rightarrow c$	7
d	$a \rightarrow d$	6
e	$a \rightarrow b \rightarrow c \rightarrow g \rightarrow e$	9
f	$a \rightarrow b \rightarrow f$	5
g	$a \rightarrow b \rightarrow c \rightarrow g$	8
h	$a \rightarrow b \rightarrow c \rightarrow g \rightarrow h$	10

Problem 4.

- a) The algorithm has two functions; MysteryWrapper(MW) function contains MysteryRecursive(MR) that is called on the first coin in $R[0]$, which is put in the set named visitedCoins. In MR, the first coin traverses to the first transaction in its associated linked list in L. Like the first coin, the transaction is put in the set named visitedTrans. Once traversed to the transaction, MR is called on the next unvisited coin in

its linked list within R. The steps above are repeated until the algorithm arrives at one of the 3 end states.

- a coin connected to a single transaction
- a transaction connected to a single visited coin (the one it traversed from)
- a transaction connected to multiple visited coins

In the first end state, the coin returns the control over to the transaction that called a recursive function on it. The coin is then mapped to the transaction in M. This process is repeated up to the very first coin. Once the control is passed over to the first coin, it finishes iterating its associated transactions and finds its pair, if any. Additionally, throughout the process, whether the graph contains any cycle is checked. If it turns out cyclic, MW clears the contents of the map and returns an empty map. This cyclic check is incorporated for all the states. In the second end state, the final transaction is mapped to the coin (within the same MR) before returning the control to the previous transaction that called MR on the coin. In the third end state, it is important to check whether each of the coins, aside from the one traversed from, connected to the final transaction is connected to a unique transaction that has only one coin in its associated linked list in R. If not, the graph is deemed cyclic. If the coin cannot sequentially travel to the last coin or the last transaction because of disconnection in the graph, MW first checks whether the disconnected portion of the graph just explored is cyclic. If it is acyclic, another MR is called on the next unvisited coin.

ALGORITHM *MysteryWrapper*(n, m, L, R)

// Input: identical to the required input from the specification

// Output: unique mapping M

DECLARE visitedCoins as set

DECLARE visitedTrans as set

SET isCyclic and endsWithTran to false

firstCoin <- the head of the first linkedlist in R

// 0 is a null value for lastTran denoting the last transaction that called MR

lastTran <- 0

INSERT firstCoin in visitedCoins

WHILE isCyclic is false and visitedCoins size is not the same as n

// L, R, visitedCoins, visitedTrans, M, isCyclic, and endsWithTran are passed as references

MysteryRecursive(firstCoin, lastTran, L, R, visitedCoins, visitedTrans, M, isCyclic, endsWithTran)

IF isCyclic is true,

CLEAR the contents of M

RETURN M

ALGORITHM

MysteryRecursive(currCoin, lastTran, L, R, visitedCoins, visitedTrans, M, isCyclic, endsWithTran)

// Output: M with coins as keys and transactions as values

currTrans <- L[currCoin]'s linked list

FOR each currTran in currTrans

IF currTran is not in visitedTrans

INSERT currTran in visitedTrans

INCREMENT numTransTravelled

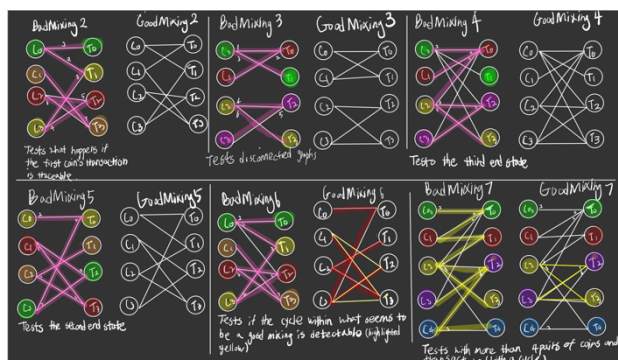
nextCoins <- R[currTran]'s linked list

FOR each nextCoin in nextCoins

```

IF nextCoin is not in visitedCoins
  INSERT nextCoin in visitedCoins
  MR(nextCoin,currTran,L,R,visitedCoins,visitedTrans,M,isCyclic,endsWithTran)
  IF isCyclic is false
    IF endsWithTran is true
      INSERT currCoin and currTran in M
    ELSE
      INSERT nextCoin and currTran in M
  IF no nextCoin is travelled // the traversal ends at a transaction
    IF R[currTran] is connected to multiple coins
      SET isAllUniquelyMapped to true
    FOR each nextCoin in nextCoins
      IF nextCoin is not currCoin
        SET isUniquelyMapped to true
        nextTrans <- L[nextCoin]
        FOR nextTran in nextTrans
          IF R[nextTran] contains only one coin
            SET isUniquelyMapped to true
          IF isUniquelyMapped is false
            SET isAllUniquelyMapped false
    IF isAllUniquelyMapped
      INSERT currCoin and currTran in M
    ELSE
      SET isCyclic to true
  ELSE // If R[currTran] is only connected to currCoin
    SET endsWithTran to true
    INSERT currCoin and currTran in M
ELSE // When exiting the recursion, check whether the graph is cyclic
  IF currTran is not lastTran // if currTran is not same as the tran that called recursion
    FOR each coin in R[currTran]
      IF coin is not currCoin
        FOR each tran in L[coin]
          IF M already has currCoin and tran mapped
            SET isCyclic to true
IF there is no unvisited currTran for currCoin
  Set endsWithTran to False

```



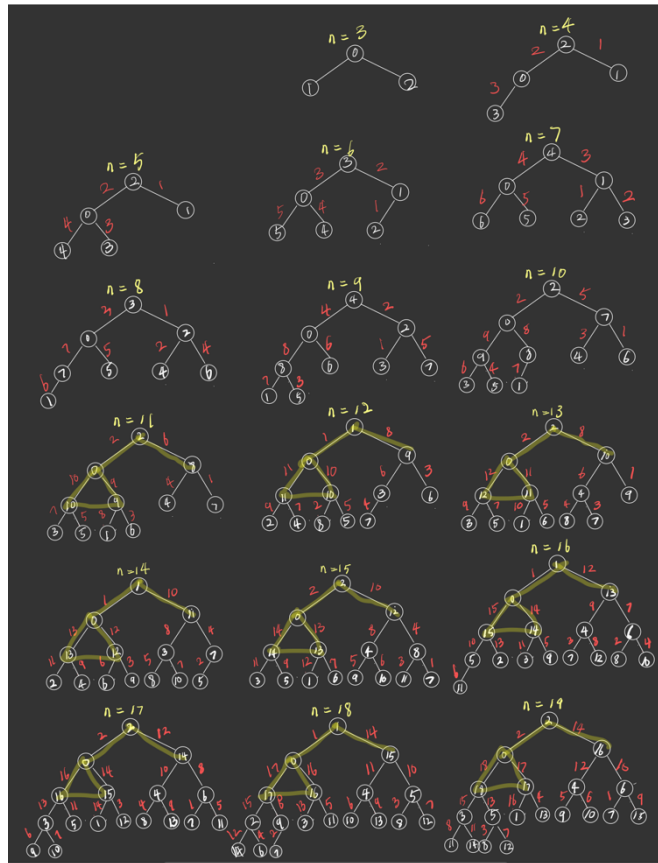
c) The left diagram shows some of the scenarios considered for the algorithm to be as error-free as possible. Unit tests for the scenarios have been written in C++ and been attached as an extra reference to demonstrate the algorithm.

d) The algorithm is guaranteed to check at worst $2n + m$ as it traverses to the end state. When it arrives at

the end state, the additional check may incur additional traversal of at most $(n + m)$. The

more edges the graph has, the earlier it will be deemed cyclic, prompting the early termination of the algorithm. Therefore, the time complexity is at most $3n + 2m$.

Problem 5.



a) The brute-force search works relatively efficiently compared to a human's manual calculation for a tree up to a certain size. However, the number of permutations factorially increases as n gets bigger, making it impossible to produce results within the required running time of 5 minutes. However, given the hard-earned results, a pattern to exploit emerges for a tree of size $n > 10$. First, the largest number $n - 1$ must be adjacent to 0 to produce the largest label, $n - 1$. Also, it seems intuitively easier to pair the second largest with 0 to produce the second largest label. As a result, the triangular positioning on the left side of the root node effectively reduces the size of the problem by 3. To compute a permutation for a tree of size 19 in time, the problem still needs to be reduced at least by 2. For bigger trees with the even number of nodes, the

root's value oddly turned out to be 1 many times; with the odd number of nodes, it evenly turned out to be 2. Finally, the right child of the root persistently contained a value of $n - 3$. Given this pattern, the maximum n to solve becomes 14, making it likely to compute a permutation for each of the trees within the required time (the pattern is highlighted in yellow in the diagram above).

As for the brute-force algorithm using recursion that is executed on trees of all sizes, it first tries a number within available numbers to see whether the leftmost element (node) whose value is None can contain the number while preserving the magical labelling property. It does so by checking whether the resultant label(s) is a duplicate already claimed by other nodes. If not, it recursively repeats the same process for the next leftmost element whose value is None. If no magical label is available, the control of the flow is returned where the recursion was last called and another num in nums is tried to try the same index yet again. This process is repeated until there are no magical labels left to choose from, in which case, all the recursion calls come off the stack, stop searches and, finally, return the permutation.

ALGORITHM *build_magical_tree*(n)

// Input: n (number of nodes in a tree)

// Output: prints the nodes of a tree

```
SET p to a list to hold the permutation
SET nums to a set of available numbers to permutate
SET mls to a set of available magical labels which is the same as nums except for the zero
IF n is equal to or less than 10
  ret <- build_magical_tree_recursive(p, 0, nums, mls, false)
  OUTPUT ret
ELSE // if n is greater than 10
  SET the elements of indices from 0 to 4 according to the pattern discussed above
  REMOVE the values used above from nums
  REMOVE the values used above from mls save for 0.
  ret <- build_magical_tree_recursive(p, 5, nums, mls, False) // start from the fifth index
  OUTPUT ret
```

ALGORITHM *build_magical_tree_recursive(p, idx, nums, mls, is_ml_found)*

// Input: p, a list to contain a valid permutation; idx, the index of a node whose value is being determined; nums, a list of available numbers; mls, a list of available magical labels; is_ml_found, a Boolean to help terminate the search once the valid permutation found
// Output: ret, the permutation that holds the magical labelling property.

```
SET ret to nothing
FOR each num in nums
  IF the final permutation is not found
    mls_copy <- updated_mls(p, idx, num, mls)
  IF mls_copy is a copy of mls:
    MAKE a copy of p, nums & SET p_copy[idx] to num & REMOVE num from nums_copy
    IF p_copy has an element (node) whose value needs to be assigned
      SET left_most_none_idx to the index of such an element
      ret <- build_magical_tree_recursive(p_copy, left_most_none_idx, nums_copy,
        mls_copy, is_ml_found)
      IF ret contains the required number of numbers
        SET is_ml_found to true
  IF there is no more mls to choose from
    SET is_ml_found to true
  ret <- p_copy
RETURN ret
```

ALGORITHM *update_mls(p, idx, num, mls)*

// This auxiliary algorithm helps determine whether the selected number results in a duplicated magical label. If not, return the updated_mls. Otherwise return nothing. Please refer to the code for more details.